

CSE 231: OS Assignment 2

Writing Your Own System Call

1 Description and implementation of the code

- SYSCALL_DEFINE2 was used as 2 input parameters were passed
- The two parameters were the pid of the required function and the name (along with the required path) of the required output file name.
- pid_task() was used to find the task_struct of the corresponding pid.
- Using the found task_struct, the required fields were printed.
- The printed fields include:
 - pid
 - task name
 - vruntime of the task
 - priority
 - normal priority
 - static priority
- A file was created using the path taken as input.
- filp_open() function was used to do the same.
- If the file with the current name is not present, a new file would be created, else the required data would be written over in the file.
- To do this, all the data would be first converted to string (character array)
- If the process is successful, the syscall would return 0 else 1.

```

SYSCALL_DEFINE2(sh_task_info, int, pid, char *, fname)
{
    struct task_struct *task;
    task = pid_task(find_vpid(pid), PIDTYPE_PID);
    if (task == NULL) {
        printk("Process with PID: %d not found\n", pid);
        return 1;
    } else {
        printk("PID: %d\n", pid);
        printk("Process: %s\n", task->comm);
        printk("vruntime: %lld\n", (task->se).vruntime);
        printk("Priority: %d\n", task->prio);
        printk("Static Priority: %d\n", task->static_prio);
        printk("Normal Priority: %d\n", task->normal_prio);
        char buf[256];
        char buf2[512] = "";
        strcat(buf2, "PID: ");
        loff_t pos;
        struct file *fp;
        mm_segment_t fs;
        long ch = strncpy_from_user(buf, fname, sizeof(buf));
        if (ch < 0 || ch == sizeof(buf)) {
            return -EFAULT;
        }
    }
}

```

```

fp = filp_open(buf, O_RDWR | O_CREAT, 0644);
if (IS_ERR(fp)) {
    int err = 0;
    err = PTR_ERR(fp);
    printk("Error Code: %d\n", err);
    return 1;
}
fs = get_fs();
set_fs(KERNEL_DS);
pos = 0;
char str1[20];
sprintf(str1, "%d", pid);
strcat(buf2, str1);
strcat(buf2, " is the process ");
strcat(buf2, task->comm);
strcat(buf2, "\n");
strcat(buf2, "vruntime is ");
char str2[20];
sprintf(str2, "%lld", (task->se).vruntime);
strcat(buf2, str2);
strcat(buf2, "\n");
strcat(buf2, "Priority is ");
char str3[20];
sprintf(str3, "%d", task->prio);
strcat(buf2, str3);

```

```

strcat(buf2, " is the process ");
strcat(buf2, task->comm);
strcat(buf2, "\n");
strcat(buf2, "vruntime is ");
char str2[20];
sprintf(str2, "%lld", (task->se).vruntime);
strcat(buf2, str2);
strcat(buf2, "\n");
strcat(buf2, "Priority is ");
char str3[20];
sprintf(str3, "%d", task->prio);
strcat(buf2, str3);
strcat(buf2, "\n");
strcat(buf2, "Static Priority is ");
char str4[20];
sprintf(str4, "%d", task->static_prio);
strcat(buf2, str4);
strcat(buf2, "\n");
strcat(buf2, "Normal Priority is ");
char str5[20];
sprintf(str5, "%d", task->normal_prio);
strcat(buf2, str5);
vfs_write(fp, buf2, strlen(buf2), &pos);
fclose(fp);
set fs(fs);

```

- The above are some snippets of the code written for the system call.

2 Inputs from the user

- The user is required to enter 2 inputs:
 - pid of the required process
 - the name (along with the address as to where to store) of the file to be created or written into while executing the system call.

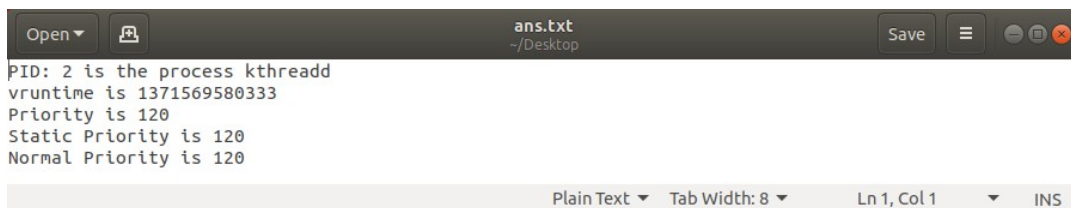
```

pw@ubuntu:~/Desktop/Sem3/OS/Assignment2/2.2$ make run
gcc test.c
./a.out
Please enter the required PID: 2
Please enter the name of the new file required along with path: /home/pw/Desktop/ans.txt
System call sys_petime returned 0

```

3 Outputs

- If the inputs are correct, the file containing the data is created.



The screenshot shows a text editor window titled 'ans.txt' with the following content:

```

PID: 2 is the process kthreadd
vruntime is 1371569580333
Priority is 120
Static Priority is 120
Normal Priority is 120

```

The editor interface includes a toolbar with 'Open', 'Save', and other icons, and a status bar at the bottom indicating 'Plain Text', 'Tab Width: 8', 'Ln 1, Col 1', and 'INS'.

- The data is also printed in the kernel logs and it can be displayed using the command `dmesg | tail`.

```

pw@ubuntu:~/Desktop/Sem3/OS/Assignment2/2.2$ dmesg | tail
[15081.684954] vruntime: 1371569580333
[15081.684954] Priority: 120
[15081.684955] Static Priority: 120
[15081.684956] Normal Priority: 120
[15424.776489] PID: 1
[15424.776493] Process: systemd
[15424.776494] vruntime: 2070883175
[15424.776495] Priority: 120
[15424.776496] Static Priority: 120
[15424.776497] Normal Priority: 120

```

- In case of an error like invalid pid error message is stored in the kernel log and file is not created.

4 Error Values

- If everything runs fine, the syscall returns 0.
- If there is an error, 1 is returned.
- Error handling has been done in both the syscall implementation.
- If a wrong (invalid) pid is entered, i.e. a process with that pid does not exist then we get a corresponding message. In such a situation, the file is not created.
- Similarly, other errors like invalid path name, invalid file size (it has been assumed that the size of file including the path can not exceed 256 characters) etc. also give their corresponding error messages.
- The above error messages can be viewed by typing the command `dmesg | tail`.

5 Diff File

- Three different diff files are supplied with this code.
- The diff files were created after executing the `make clean` command.
- The diff files by the name of “diff.txt” and “diffruN.txt” were created by using the flags “-r”, “-u” and “-N”.
- The diff file by the name of “diffruaN.txt” was created by using the flags “-r”, “-u”, “-a” and “-N”.

